

Review

Homework 1

PEAS: Performance Measure, Environment, Actuators, Sensors

- **Performance Measure** – Criteria for evaluating the agent's success.
- **Environment** – The surroundings in which the agent operates.
- **Actuators** – Components that allow the agent to take actions.
- **Sensors** – Components that enable the agent to perceive its environment.

Formulating Search Problems

1. **Problem:** Slide tiles on a 3x3 grid to match a target configuration.
2. **State Space:** All possible arrangements of 8 numbers and a blank tile (e.g., "1 2 3, 4 5 6, 7 8 _").
3. **Initial State:** A specific scrambled board (e.g., "1 2 3, 4 8 5, 7 6 _").
4. **Goal State:** The solved board (e.g., "1 2 3, 4 5 6, 7 8 _").
5. **Actions:** Move the blank tile up, down, left, or right (if legal).
6. **Transition Model:** If blank is at (2,2), "move up" → blank swaps with (1,2).
7. **Cost Function:** Each move costs 1 (or skip for BFS/DFS, since cost isn't considered).

Search Trees

Generation

- Avoid duplicate choices at the same level.
- Search trees can't have cycles but may be infinite if the state-space graph has cycles.
- A finite, acyclic state-space graph ensures a finite search tree.

GRAPH Search

Each of these algorithms create a search tree from the given state-space graph.

Tie-breakers needed to avoid multiple solutions in terms of state-expansion order.

BFS

- **Queue-based** traversal.
- **No heuristics** unless explicitly stated.
- **Stops when the goal is expanded.**

DFS

- **Stack-based** traversal.
- **No heuristics** unless explicitly stated.
- **Stops when the goal is expanded.**

UCS (Uniform Cost Search)

- **No heuristics** unless explicitly stated.
- **No revisits:** Use a **visited list** to track explored nodes.
- **Track total cost:** Sum the path costs from the start state to each node.
 - Example: If $S \rightarrow A = 5$ and $A \rightarrow B = 6$, then $S \rightarrow A \rightarrow B = 11$. Write 5 at A and 11 at B.
- **Expand the cheapest node** on the **frontier** (unexpanded nodes).
- **Stops when the goal state is expanded.**
- **Guaranteed minimum-cost path** if costs are **non-negative**.

A*

- **Uses heuristics.**
- **Admissibility** ensures the shortest path in graphs with **non-negative edge costs**.
- **Consistency** guarantees **admissibility** and prevents **revisiting/re-expanding**, ensuring **optimality**.
- **Extends UCS** by considering both real cost and heuristic cost:

$$f(n) = g(n) + h(n)$$

- **Tracking visited nodes:** Store the A* score $f(n)$.
 - In UCS, visited nodes could be ignored.
 - In A*, however, revisit a node **only if** a lower $f(n)$ score is found.

Definitions

- **Admissibility:** A heuristic is **admissible** if it **never overestimates** the cost to the goal. That is, the true cost is always **at least** the heuristic:

$$h(n) \leq h^*(n)$$

where $h^*(n)$ is the actual cheapest cost from n to G .

- **Consistency:** A heuristic is **consistent** if it satisfies the **triangle inequality**:

$$h(n) \leq h(m) + c(n, m)$$

where $c(n, m)$ is the cost of moving from n to m .

- A **consistent** heuristic is always **admissible**, but an **admissible** heuristic is not necessarily **consistent**.
- **State space** is the set of all possible states the agent can reach in the grid. Its size depends on what information is included in *each* state.
 - For a grid with $M \times N$ cells, if the state is defined by the agent's **position** (x, y) and **facing direction** (left, right, up, or down), the total number of states is:

$$M \times N \times 4$$

- There are $M \times N$ states for each direction (left, right, up, and down).
- Since there are 4 possible directions, the total state space is:

$$(M \times N) + (M \times N) + (M \times N) + (M \times N) = M \times N \times 4$$

Homework 2

Minimax with Alpha-Beta Pruning

- **Alpha:** Best explored value for the maximizer.
- **Beta:** Best explored value for the minimizer.
- Maximizer nodes update **alpha**.
- Minimizer nodes update **beta**.

Initial Values

Minimizer and maximizer start with worst-case values:

- **Maximizer:** $v = -\infty$, so **alpha** = $-\infty$.
- **Minimizer:** $v = \infty$, so **beta** = $+\infty$.
- If **alpha** is unknown, assume $-\infty$.
- If **beta** is unknown, assume $+\infty$.

Recursive Formula

- Set v to the parent node's value (or its initial default).
- **Minimizer**: If $v' < v$, update v with v' .
- **Maximizer**: If $v' > v$, update v with v' .
- **Pruning**:
 - If the node is a **maximizer** and $v' > \text{beta}$ → **prune**.
 - If the node is a **minimizer** and $v' < \text{alpha}$ → **prune**.

Difference Between Expectiminimax and Minimax

Feature	Minimax	Expectiminimax
Used In	Deterministic games (e.g., Chess, Tic-Tac-Toe)	Stochastic (randomized) games (e.g., Backgammon, Monopoly)
Types of Nodes	Max (agent's turn) and Min (opponent's turn)	Max, Min, and Chance (random events)
Opponent Modeling	Assumes opponent plays optimally	Accounts for randomness using probability
Decision Process	Chooses the move that maximizes the worst-case outcome	Computes an expected value for chance nodes
Tree Structure	Alternates between Max and Min levels	Alternates between Max, Min, and Chance levels
Evaluation	Uses a utility function directly	Uses a weighted average of possible outcomes at chance nodes

For expectiminimax:

$$\text{Weighted Probability} = \frac{\sum_{i=1}^n P(n_i) \cdot V(n_i)}{\text{Number of Chance Nodes}}$$

Homework 3

Formulas

$$V_k^{\pi_i}(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{k-1}^{\pi_i}(s')]$$

Here are the fundamental equations for **Q-learning**:

Q-Value Update Equation

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

Where:

- $Q(s, a)$ is the current Q-value for state s and action a .
- α is the learning rate ($0 < \alpha \leq 1$).
- r is the reward received after taking action a in state s .
- γ is the discount factor ($0 \leq \gamma \leq 1$).
- s' is the next state after taking action a .
- $\max_{a'} Q(s', a')$ is the highest Q-value of the next state.

Greedy Action Selection (Exploitation)

$$a^* = \arg \max_a Q(s, a)$$

This selects the action a^* with the highest Q-value for state s .

Epsilon-Greedy Policy (Exploration vs. Exploitation)

With probability ϵ , choose a random action; otherwise, choose the best action:

$$a = \begin{cases} \text{random action,} & \text{with probability } \epsilon \\ \arg \max_a Q(s, a), & \text{with probability } 1 - \epsilon \end{cases}$$

Feature-based Representation

Q-learning updates the weights of features rather than maintaining a table of $Q(s, a)$ values for each state-action pair. This is useful when the state space is large or continuous.

Q-Value Approximation

Instead of using a table, the Q-value is represented as a **linear function** of features:

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \cdots + w_n f_n(s, a) = \sum_i w_i f_i(s, a)$$

where:

- $f_i(s, a)$ are the features describing the state-action pair.
- w_i are the corresponding weights.

Weight Update Rule

The weights are updated using a gradient descent step:

$$w_i \leftarrow w_i + \alpha \delta f_i(s, a)$$

where:

- α is the learning rate.
- δ is the **TD error**, given by:

$$\delta = \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$